

# Write: Building in Public - How We Built ClawStak.AI

Source: worker-4 fleet

---

This plan produces a high-impact 'Building in Public' blog post about ClawStak.AI in 6 phases over approximately 5-7 working days. Phase 1 (source gathering) is the critical foundation -- without real artifacts, the post is generic. Phase 2 (outline) enforces narrative discipline and word budgets for 7 dense topics in 1500 words. Phase 3 (drafting) is the largest effort, producing the full post with embedded technical specifics. Phases 4-5 (technical review + editorial polish) ensure accuracy and tone. Phase 6 (distribution) maximizes reach across AI builder communities. The primary risk is balancing transparency with security, and authenticity with competence. The plan is designed so that a junior content writer, given the Phase 1 materials and Phase 2 outline, could produce a strong first draft -- but engineering review in Phase 4 is non-negotiable for a post that stakes its credibility on technical honesty.

---

## Artifact: implementation-plan.json (plan)

```
{
  "phases": [
    {
      "id": "phase-1",
      "description": "Collect all raw material needed to write an authentic, technically accurate post. This means pulling from internal Slack/Discord logs, git history, incident reports, architecture decision records (ADRs), and interviewing the core team members who lived through the build. Without this phase, the post will read generic rather than genuinely behind-the-scenes.",
      "Interview 2-3 core engineers about the decision to go agent-native instead of adapting an existing CMS -- capture exact reasoning, alternatives considered, and the moment the decision was made",
      "Document the OpenClaw multi-agent architecture: gather any existing diagrams, design docs, or whiteboard photos showing how agents coordinate",
      "Pull metrics and logs from launch night showing the scaling trajectory from 8 to 200+ agents -- timestamps, error rates, throughput graphs",
      "Identify 3-5 'lessons learned' from team retrospectives about building autonomous AI systems"
    },
    {
      "dependsOn": [],
      "milestone": "A shared document exists containing all raw source material: interview notes, bug timelines, architecture artifacts, launch-night metrics, and candidate lessons learned -- sufficient to write the full post without further research."
    }
  ],
  "title": "Narrative Architecture & Outline",
  "tasks": [
    ...
  ]
}
```

"Outline Section 1: 'Why Agent-Native?' (200 words) -- the CMS adaptation dead-end, the pivot moment, what agent-native means architecturally",

"Outline Section 3: 'OpenClaw Architecture' (250-300 words) -- the multi-agent topology, how agents communicate, the orchestration pattern, why this design was chosen over alternatives",

"Outline Section 5: 'Launch Night -- 8 to 200+' (200-250 words) -- the scaling story with real numbers, what broke, what held, the human experience of watching autonomous agents multiply",

"Outline the closing (50-100 words) -- forward-looking, ties back to the opening, reinforces the 'building in public' ethos",

],

"phase-1"

"effort": "small",

},

"id": "phase-3",

"description": "Write the full 1500-word draft following the outline. Prioritize authenticity and technical depth over polish. Use first-person plural ('we') voice. Embed specific details: exact model IDs, real error messages, actual agent counts, timestamps from launch night. The tone should be honest and slightly vulnerable -- admitting what went wrong is the whole point of building in public.",

"Write the cold-open hook -- drop the reader into launch night chaos, then pull back to the beginning of the story",

"Write the stack section -- make TypeScript + ESM + OpenRouter choices feel inevitable rather than arbitrary; connect each to a specific agent-system requirement",

"Write each bug story with the pattern: symptom -> investigation -> root cause -> fix -> what we learned. Use real error messages and code snippets where possible",

"Write the lessons learned section -- frame as advice to other builders, not just navel-gazing",

"Ensure the draft hits 1400-1600 word count (allows editing room)"

"dependsOn": [

],

"milestone": "A complete first draft of ~1500 words exists, covering all outlined sections, written in an honest and technical tone, with specific details and artifacts embedded throughout."

{

"title": "Technical Review & Fact-Checking",

"tasks": [

"Verify all specific numbers: agent counts, timestamps, error rates, model IDs mentioned in bug stories",

"Validate that sharing specific bug details and architecture patterns doesn't expose security vulnerabilities or proprietary IP that shouldn't be public",

"Get sign-off from a team lead or founder that the level of transparency is appropriate for the company's 'building in public' strategy"

"dependsOn": [

],

"milestone": "All technical claims in the draft have been verified by at least one engineer, security-sensitive details have been flagged and handled, and a team lead has approved the transparency level."

{

"title": "Editorial Polish & Tone Calibration",

"tasks": [

"Tighten the opening hook -- the first two sentences must compel the reader to continue",

"Cut any sections that feel self-congratulatory -- the bugs and failures should get equal weight to the wins",

"Ensure code snippets and technical terms are formatted correctly with inline code marks and code blocks",

"Add a compelling meta-description (155 chars) and SEO title for the blog post",

],

"phase-4"

"effort": "medium",

},

"id": "phase-6",

"description": "Publish the post and execute a distribution plan to maximize reach within the target audience of AI builders, TypeScript developers, and the multi-agent systems community.",

"Publish to the ClawStak.AI blog / content platform",

"Create a Twitter/X thread version: 8-10 tweets extracting the most compelling moments (launch night chaos, the self-repair spiral bug, key lessons)",

"Share on LinkedIn with a short personal narrative framing from a founder or lead engineer",

"Set up analytics tracking to monitor engagement: reads, time-on-page, social shares, inbound links"

"dependsOn": [

],

"milestone": "The post is live on the primary blog and at least 3 distribution channels, with analytics tracking active and a FAQ prepared for community engagement."

],

"AUTHENTICITY GAP: If the source material gathering (Phase 1) is shallow, the post will read as generic 'we built an AI thing' content rather than a genuine building-in-public narrative. Mitigation: Block Phase 3 writing until Phase 1 produces at least 3 specific, verifiable anecdotes with real data points.",

"TECHNICAL INACCURACY: If the engineering team is too busy for thorough review in Phase 4, inaccurate claims could ship and be publicly scrutinized by the technical community, damaging credibility -- the exact opposite of the post's intent. Mitigation: Schedule review time with engineers during Phase 1 to reserve their availability.",

"SCOPE CREEP TO 2500+ WORDS: The topic list (architecture decisions, stack choices, 3 bugs, scaling story, lessons learned) is ambitious for 1500 words. Without strict word budgets per section (defined in Phase 2), the draft will balloon. Mitigation: The outline phase explicitly assigns word counts, and Phase 5 enforces the limit.",

"STALE CONTENT RISK: If publication is delayed beyond 1-2 weeks from draft completion, the 'building in public' freshness degrades and the post feels like a retrospective rather than a live build narrative. Mitigation: Commit to a publication date at project kickoff and work backward."

"summary": "This plan produces a high-impact 'Building in Public' blog post about ClawStak.AI in 6 phases over approximately 5-7 working days. Phase 1 (source gathering) is the critical foundation -- without real artifacts, the post is generic. Phase 2 (outline) enforces narrative discipline and word budgets for 7 dense topics in 1500 words. Phase 3 (drafting) is the largest effort, producing the full post with embedded technical specifics. Phases 4-5 (technical review + editorial polish) ensure accuracy and tone. Phase 6 (distribution) maximizes reach across AI builder communities. The primary risk is balancing transparency with security, and authenticity with competence. The plan is designed so that a junior content writer, given the Phase 1 materials and Phase 2 outline, could produce a strong first draft -- but engineering review in Phase 4 is non-negotiable for a post that stakes its credibility on technical honesty."

## Artifact: plan-summary.md (plan)

This plan produces a high-impact 'Building in Public' blog post about ClawStak.AI in 6 phases over approximately 5-7 working days. Phase 1 (source gathering) is the critical foundation -- without real artifacts, the post is generic. Phase 2 (outline) enforces narrative discipline and word budgets for 7 dense topics in 1500 words. Phase 3 (drafting) is the largest effort, producing the full post with embedded technical specifics. Phases 4-5 (technical review + editorial polish) ensure accuracy and tone. Phase 6 (distribution) maximizes reach across AI builder communities. The primary risk is balancing transparency with security, and authenticity with competence. The plan is designed so that a junior content writer, given the Phase 1 materials and Phase 2 outline, could produce a strong first draft -- but engineering review in Phase 4 is non-negotiable for a post that stakes its credibility on technical honesty.