

# Write: Why Every AI Company Needs an Agent Fleet

Source: worker-2 fleet

---

## Final Deliverable: "Why Every AI Company Needs an Agent Fleet"

### What Was Accomplished

A complete 1,850-word thought leadership article was researched, drafted, and reviewed through a three-step specialist workflow:

#### Step 1 -- Research (Complete)

Comprehensive research was conducted across multiple sources covering:

- Industry adoption data: 72% of enterprise AI projects now use multi-agent architectures; market projected at \$52B by 2030
- Cost efficiency evidence: 80%+ cost reduction through smart model routing; 45% faster resolution; 60% more accurate outcomes
- Sources included synlabs.io, secondsedge.com, launchyourclaw.com, and digitalapplied.com

#### Step 2 -- Draft (Complete)

A full article was produced covering all five required topics:

2. Power of specialization -- domain expertise, parallel execution, compounding organizational learning
4. Cost efficiency -- 80%+ savings via model routing, failure isolation preventing cascade collapse, elastic on-demand scaling

#### Step 3 -- Review (Complete)

A thorough editorial review identified actionable improvements graded at B overall, A- potential.

---

### Remaining Items Before Publication

#### Critical (Must Fix)

#	Issue	Action Required
1	Overlength -- 1,850 words vs. 1,200 target	Cut ~650 words using the section-by-section reduction table provided in the review
3	Statistics need qualification -- "77% cost reduction" and "70% token reduction" are single case studies presented as universal	Add "In one documented deployment..." or similar qualifiers

**Important (Should Fix)**

#	Issue	Action Required
---	-------	-----------------

5	Repetitive cost/specialization arguments across 4+ sections	Consolidate into single authoritative treatment per topic
---	---	---

7	Promotional tone in ClawStak sections	Soften "ClawStak.AI makes this simple" to neutral professional language
---	---------------------------------------	---

**Nice to Have**

- Add a single-agent vs. multi-agent comparison table for scannability
  - Add visual callout boxes for key statistics
- 

**Recommended Next Step**

Run one more editing pass applying the review's priority fixes -- particularly the 35% word count reduction and claim substantiation -- then the article is ready for publication. The structural outline for a tight 1,200-word version is included in the review (Step 3) and can be handed directly to a writer or editing agent.

---

**Artifact: execution-plan.md (plan)**

```
{
  "steps": [
    {
      "id": "step-1",
      "agentRole": "researcher",
      "label": "step-1"
    },
    {
      "id": "step-2",
      "agentRole": "coder",
      "label": "step-2"
    },
    {
      "id": "step-3",
      "agentRole": "reviewer",
      "label": "step-3"
    }
  ]
}
```

{

## Artifact: synthesis.md (analysis)

# Final Deliverable: "Why Every AI Company Needs an Agent Fleet"

## What Was Accomplished

A complete 1,850-word thought leadership article was researched, drafted, and reviewed through a three-step specialist workflow:

### Step 1 -- Research (Complete )

Comprehensive research was conducted across multiple sources covering:

- Industry adoption data: 72% of enterprise AI projects now use multi-agent architectures; market projected at \$52B by 2030
- Cost efficiency evidence: 80%+ cost reduction through smart model routing; 45% faster resolution; 60% more accurate outcomes
- Sources included synlabs.io, secondsedge.com, launchyourclaw.com, and digitalapplied.com

### Step 2 -- Draft (Complete )

A full article was produced covering all five required topics:

2. Power of specialization -- domain expertise, parallel execution, compounding organizational learning
4. Cost efficiency -- 80%+ savings via model routing, failure isolation preventing cascade collapse, elastic on-demand scaling

### Step 3 -- Review (Complete )

A thorough editorial review identified actionable improvements graded at B overall, A- potential.

## Remaining Items Before Publication

### Critical (Must Fix)

#	Issue	Action Required
1	Overlength -- 1,850 words vs. 1,200 target	Cut ~650 words using the section-by-section reduction table provided in the review
3	Statistics need qualification -- "77% cost reduction" and "70% token reduction" are single case studies presented as universal	Add "In one documented deployment..." or similar qualifiers

## Important (Should Fix)

| # | Issue | Action Required |

| 5 | Repetitive cost/specialization arguments across 4+ sections | Consolidate into single authoritative treatment per topic |

| 7 | Promotional tone in ClawStak sections | Soften "ClawStak.AI makes this simple" to neutral professional language |

## Nice to Have

- Add a single-agent vs. multi-agent comparison table for scannability
  - Add visual callout boxes for key statistics
- 

## Recommended Next Step

Run one more editing pass applying the review's priority fixes -- particularly the 35% word count reduction and claim substantiation -- then the article is ready for publication. The structural outline for a tight 1,200-word version is included in the review (Step 3) and can be handed directly to a writer or editing agent.